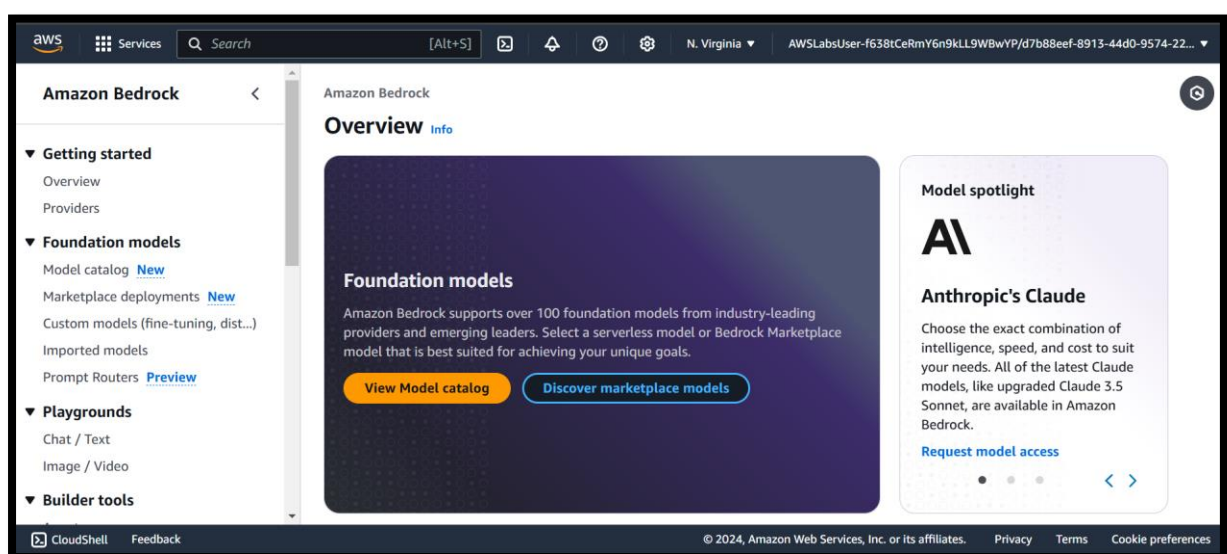
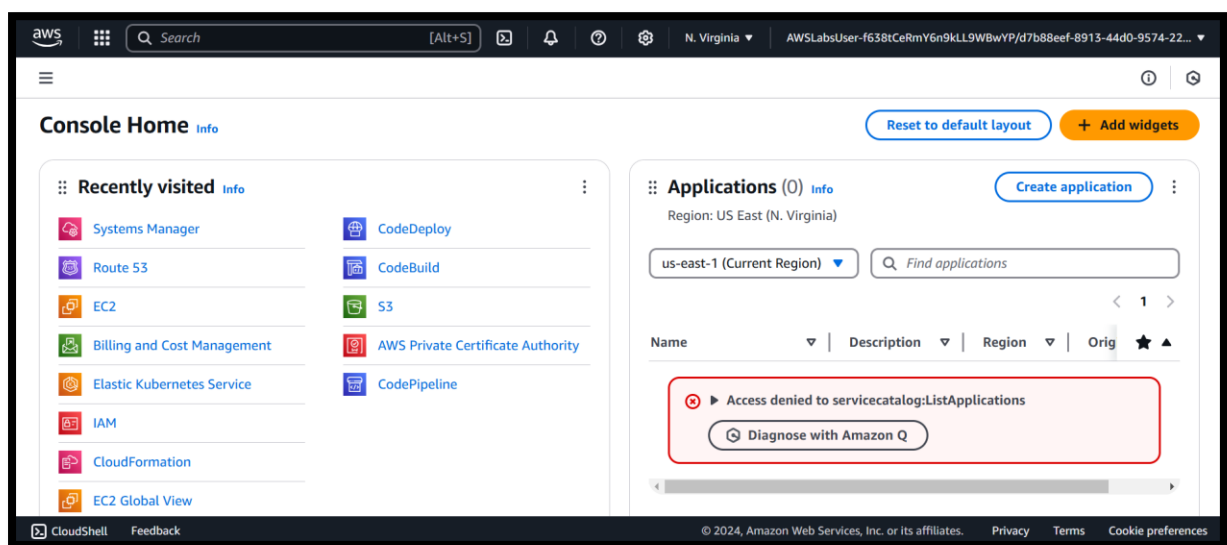


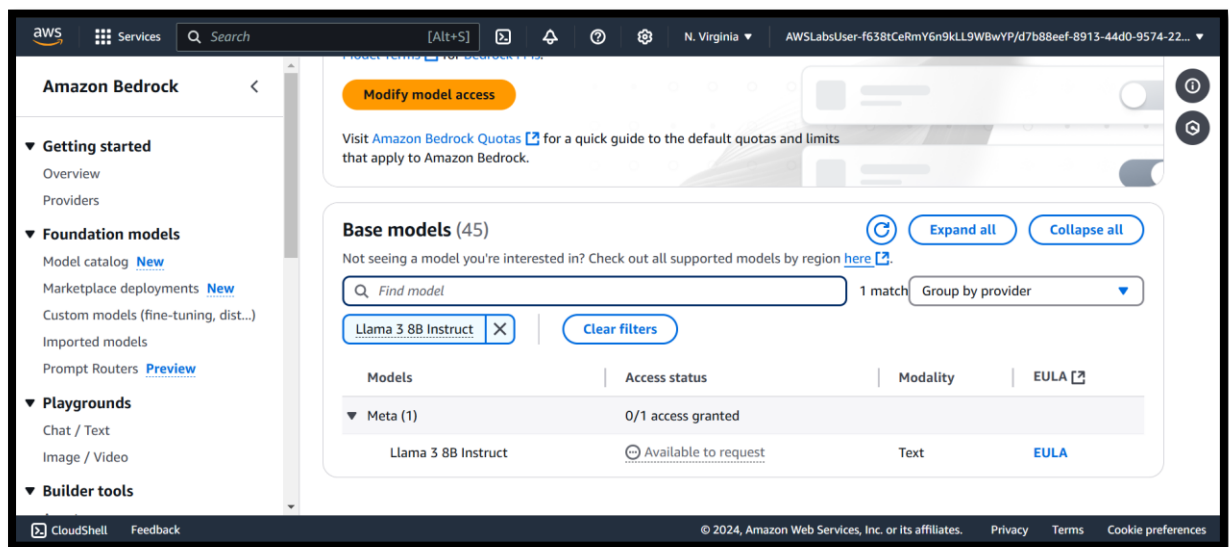
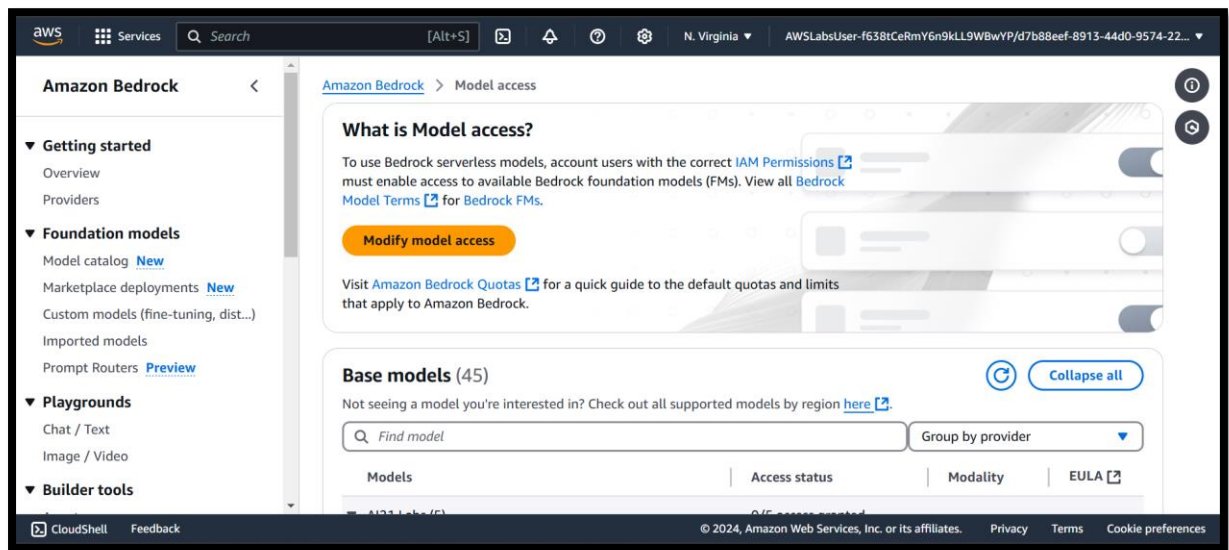
**Objective:** To run two notebook files: Task1a.ipynb, which invokes an Amazon Bedrock model for text generation using a zero-shot prompt, and Task1b.ipynb, which uses the LangChain framework to communicate with the Amazon Bedrock API and creates a custom LangChain prompt template to add context to the text generation request.

## Task 0: Set up the environment

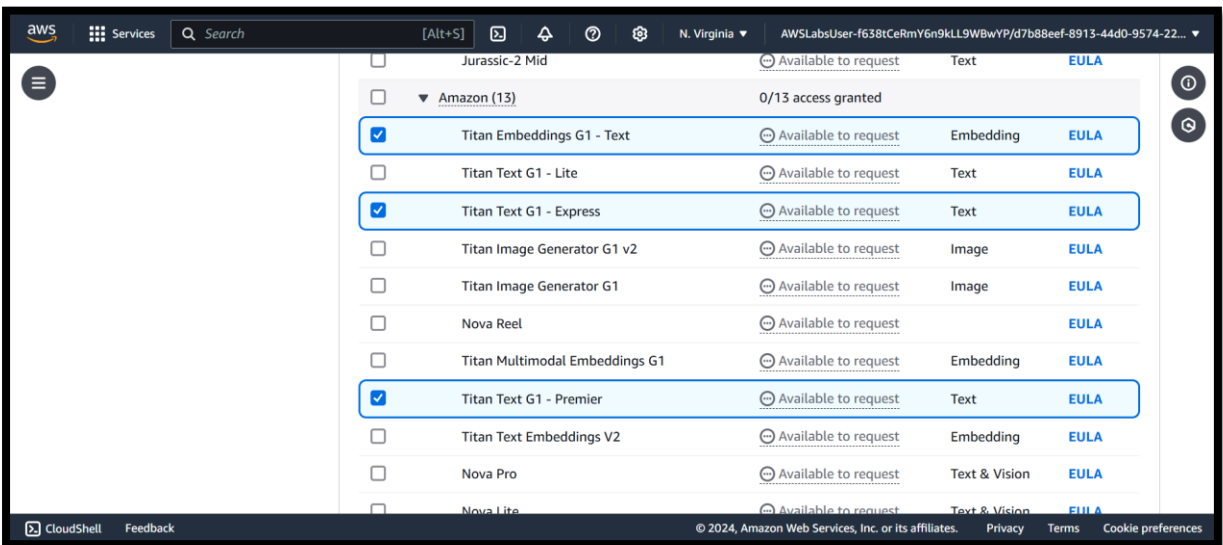
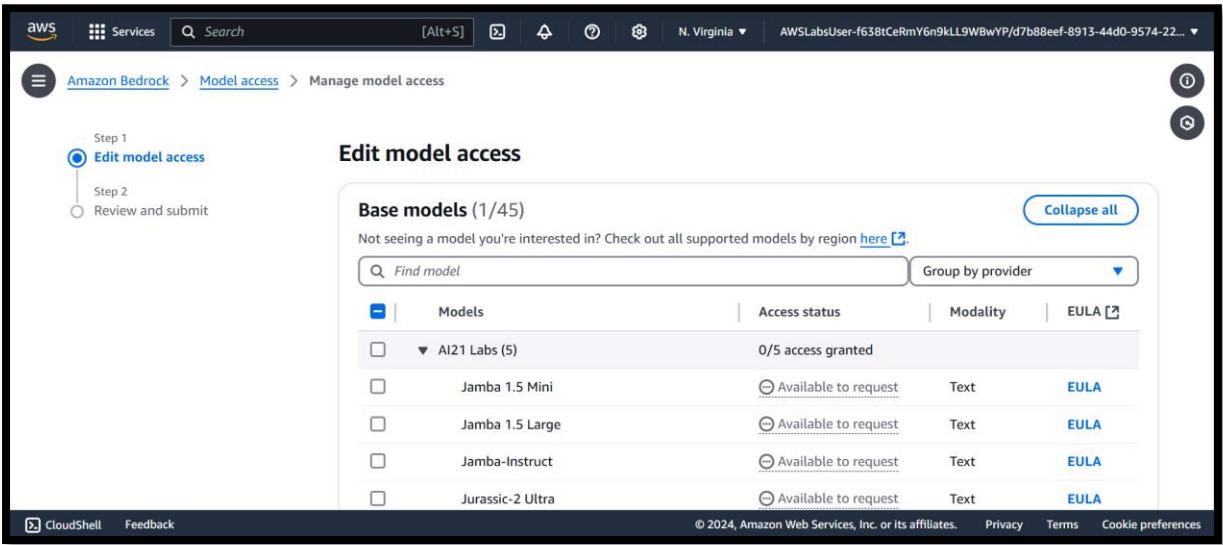
In this task, I registered the base models in the Amazon Bedrock console and launched an Amazon SageMaker Studio application to access my lab resources.

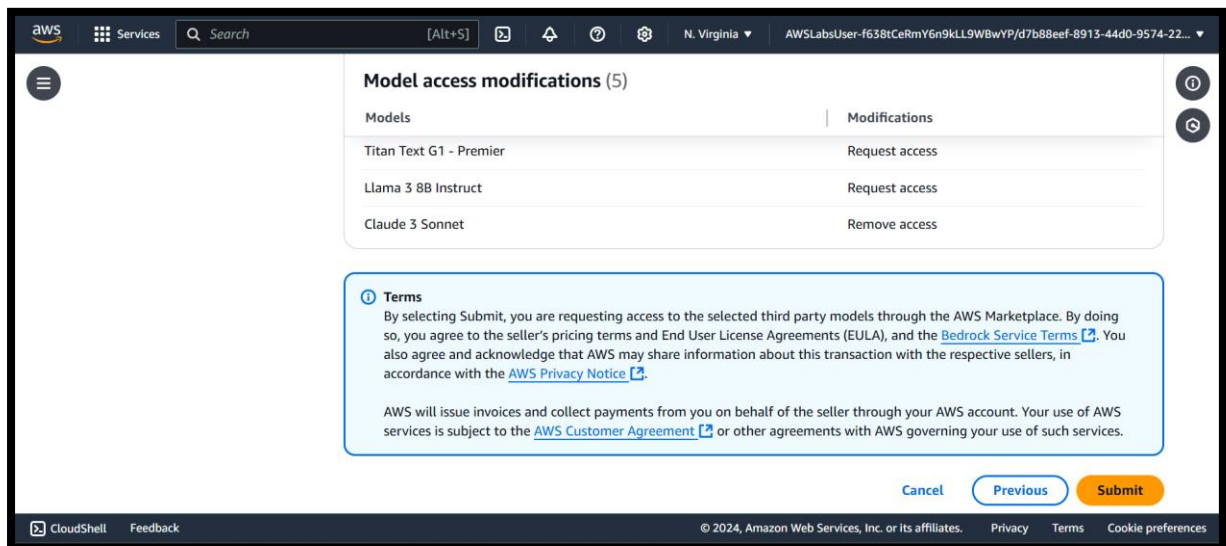
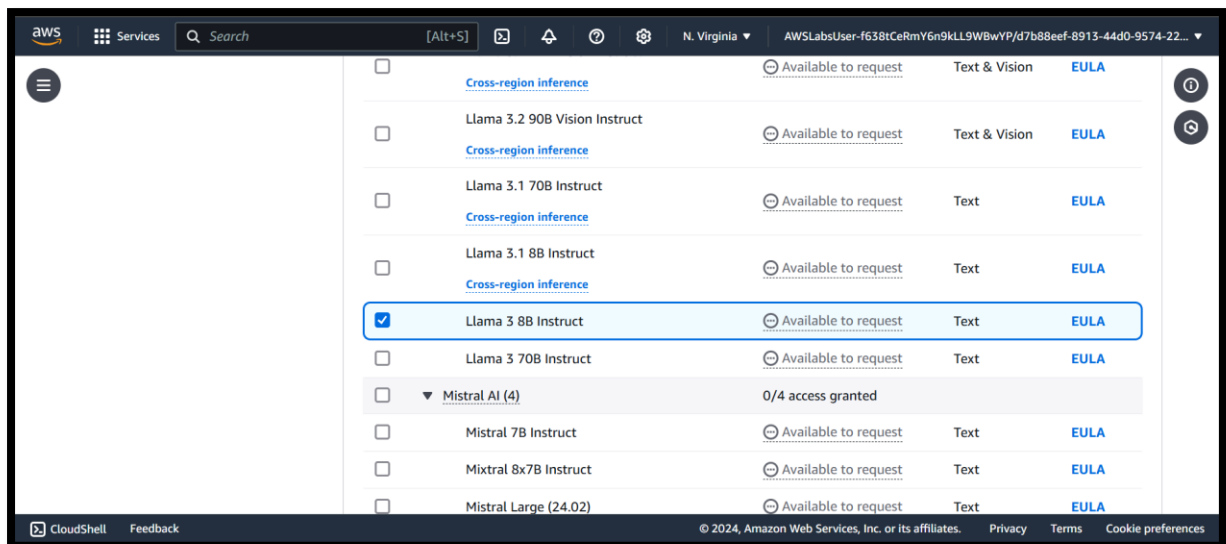


I reviewed the Access Status for each of the models. If the Access Status for one or more of the models was set to Available to request, I expanded this menu and followed the steps to enable access for them.

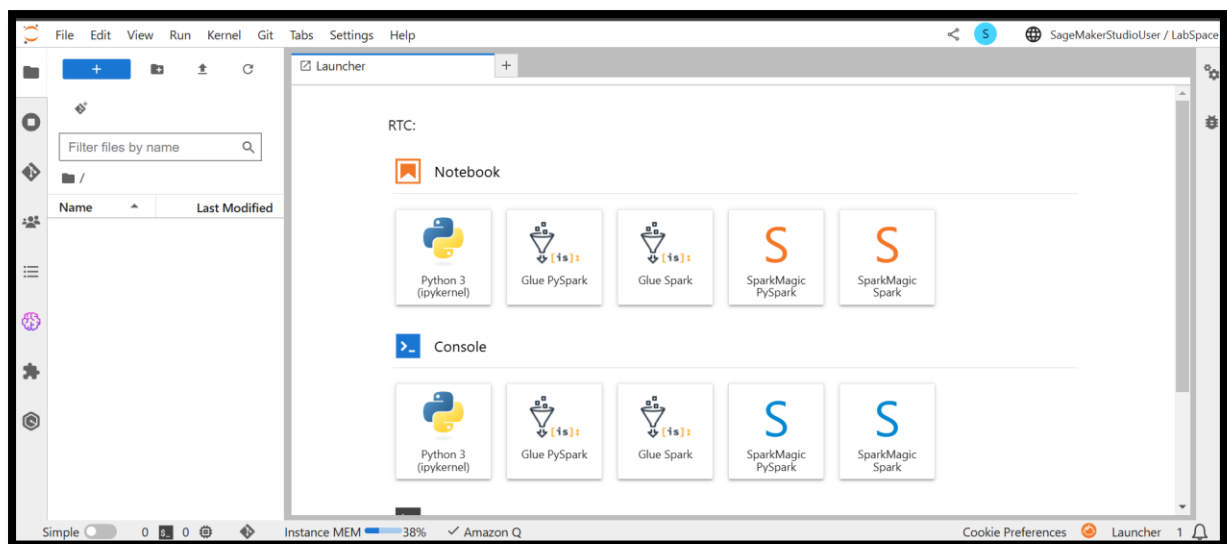


I chose Modify model access at the top of the screen.



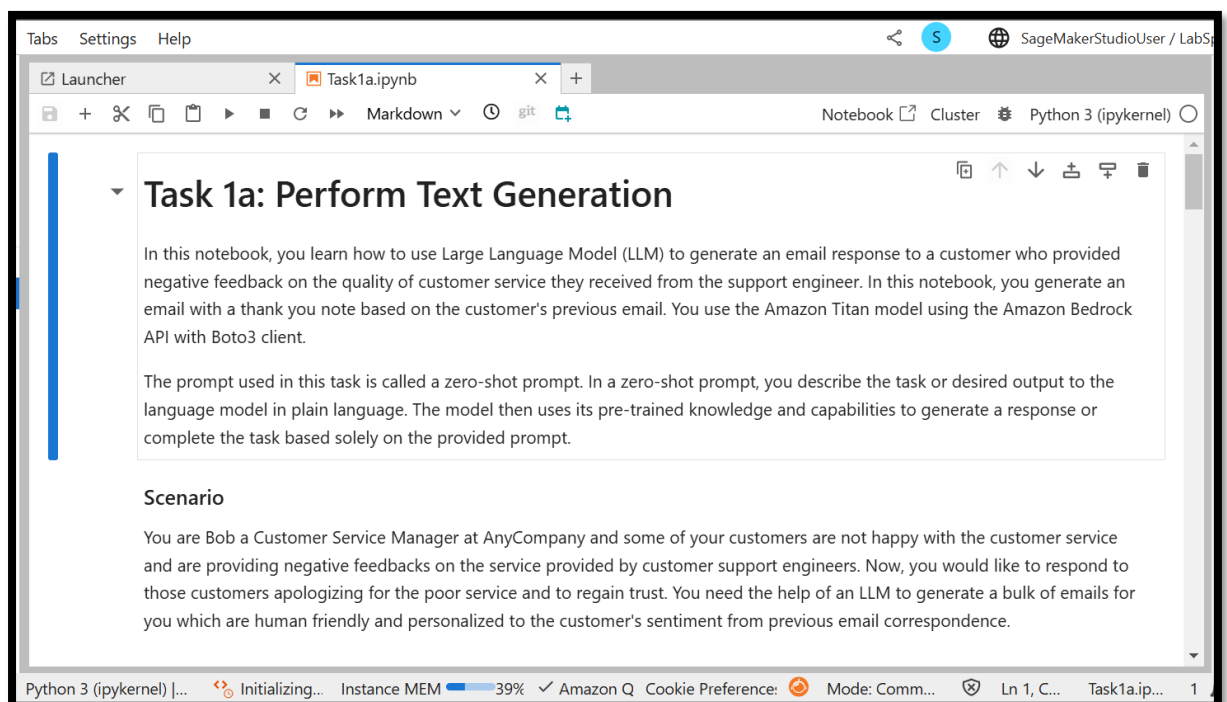


Launch an Amazon SageMaker Studio application



## Task 1: Perform Text Generation

Task1a.ipynb, which invoked an Amazon Bedrock model for text generation using a zero-shot prompt



Task1a.ipynb

Python 3 (ipykernel)

You are Bob a Customer Service Manager at AnyCompany and some of your customers are not happy with the customer service and are providing negative feedbacks on the service provided by customer support engineers. Now, you would like to respond to those customers apologizing for the poor service and to regain trust. You need the help of an LLM to generate a bulk of emails for you which are human friendly and personalized to the customer's sentiment from previous email correspondence.

## Task 1a.1: Environment setup

In this task, you set up your environment.

```
[1]: #Create a service client by name using the default session.
import json
import os
import sys

import boto3
import botocore

module_path = ".."
sys.path.append(os.path.abspath(module_path))

bedrock_client = boto3.client('bedrock-runtime', region_name=os.environ.get("AWS_DEFAULT_REGION", None))
```

Python 3 (ipykernel) Initialized (additional servers n Instance MEM 4 ✓ Amazon Cookie Preferen Mode: Com... Ln 1, ... Task1a.i... 1

Task1a.ipynb

Python 3 (ipykernel)

```
bedrock_client = boto3.client('bedrock-runtime', region_name=os.environ.get("AWS_DEFAULT_REGION", None))
```

## Task 1a.2: Generate text

In this task, you prepare an input for the Amazon Bedrock service to generate an email.

```
[2]: # create the prompt
prompt_data = """
Command: Write an email from Bob, Customer Service Manager, AnyCompany to the customer "John Doe"
who provided negative feedback on the service provided by our customer support
engineer"""

[3]: body = json.dumps({
    "inputText": prompt_data,
    "textGenerationConfig": {
        "maxTokenCount": 8192,
        "stopSequences": [],
        "temperature": 0,
        "topP": 0.9
    }
})
```

Python 3 (ipykernel) Initialized (additional servers n Instance MEM 4 ✓ Amazon Cookie Preferen Mode: Com... Ln 1, ... Task1a.i... 1

Next, you use the Amazon Titan model.

**Note:** Amazon Titan supports a context window of ~4k tokens and accepts the following parameters:

- `inputText` : Prompt to the LLM
- `textGenerationConfig` : These are the parameters that model will take into account while generating the output.

The Amazon Bedrock API provides you with an API `invoke_model` which accepts the following:

- `modelId` : This is the model ARN for the various foundation models available under Amazon Bedrock
- `accept` : The type of input request
- `contentType` : The content type of the output
- `body` : A json string consisting of the prompt and the configurations

Refer to [documentation](#) for Available text generation model Ids.

## Task 1a.3: Invoke the Amazon Titan Large language model

In this task, you explore how the model generates an output based on the prompt created earlier.

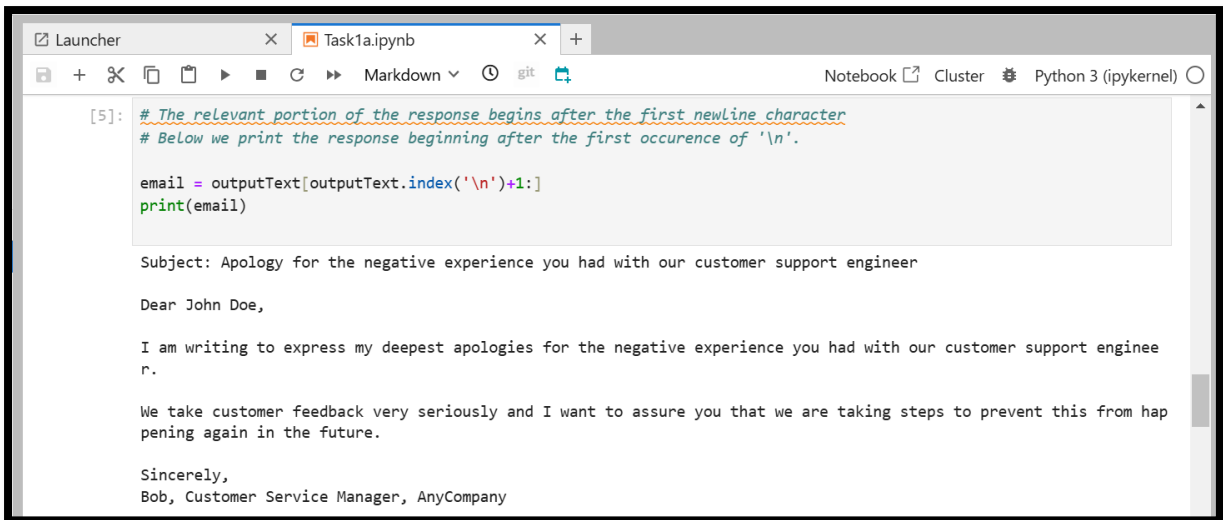
```
[4]: #invoke_model
modelId = 'amazon.titan-text-express-v1' # change this to use a different version from the model provider
accept = 'application/json'
contentType = 'application/json'
outputText = ""
try:

    response = bedrock_client.invoke_model(body=body, modelId=modelId, accept=accept, contentType=contentType)
    response_body = json.loads(response.get('body').read())

    outputText = response_body.get('results')[0].get('outputText')

except botocore.exceptions.ClientError as error:

    if error.response['Error']['Code'] == 'AccessDeniedException':
        print(f"\n{error.response['Error']['Message']}")
        \nTo troubleshoot this issue please refer to the following resources.\
        \nhttps://docs.aws.amazon.com/IAM/latest/UserGuide/troubleshoot_access-denied.html\
        \nhttps://docs.aws.amazon.com/bedrock/latest/userguide/security-iam.html\n")
    else:
        raise error
```



The screenshot shows a Jupyter Notebook interface with a tab labeled 'Task1a.ipynb'. The code cell contains a Python script that processes a response string. The output of the cell is a formatted email text.

```
[5]: # The relevant portion of the response begins after the first newline character
# Below we print the response beginning after the first occurrence of '\n'.

email = outputText[outputText.index('\n')+1:]
print(email)
```

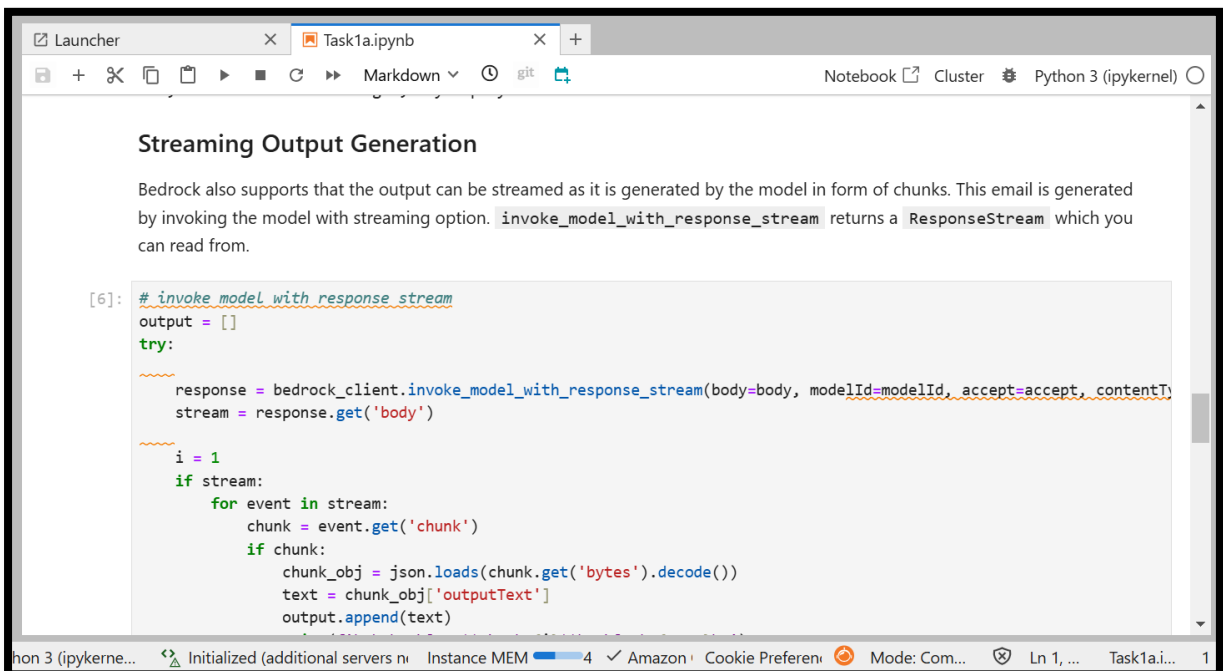
Subject: Apology for the negative experience you had with our customer support engineer

Dear John Doe,

I am writing to express my deepest apologies for the negative experience you had with our customer support engineer.

We take customer feedback very seriously and I want to assure you that we are taking steps to prevent this from happening again in the future.

Sincerely,  
Bob, Customer Service Manager, AnyCompany



The screenshot shows a Jupyter Notebook interface with a tab labeled 'Task1a.ipynb'. The code cell contains a Python script that demonstrates streaming output generation. The output of the cell is a formatted email text.

### Streaming Output Generation

Bedrock also supports that the output can be streamed as it is generated by the model in form of chunks. This email is generated by invoking the model with streaming option. `invoke_model_with_response_stream` returns a `ResponseStream` which you can read from.

```
[6]: # invoke_model_with_response_stream
output = []
try:
    response = bedrock_client.invoke_model_with_response_stream(body=body, modelId=modelId, accept=accept, contentType=contentType)
    stream = response.get('body')

    i = 1
    if stream:
        for event in stream:
            chunk = event.get('chunk')
            if chunk:
                chunk_obj = json.loads(chunk.get('bytes').decode())
                text = chunk_obj['outputText']
                output.append(text)
```

h3> Streaming Output Generation

Bedrock also supports that the output can be streamed as it is generated by the model in form of chunks. This email is generated by invoking the model with streaming option. `invoke_model_with_response_stream` returns a `ResponseStream` which you can read from.

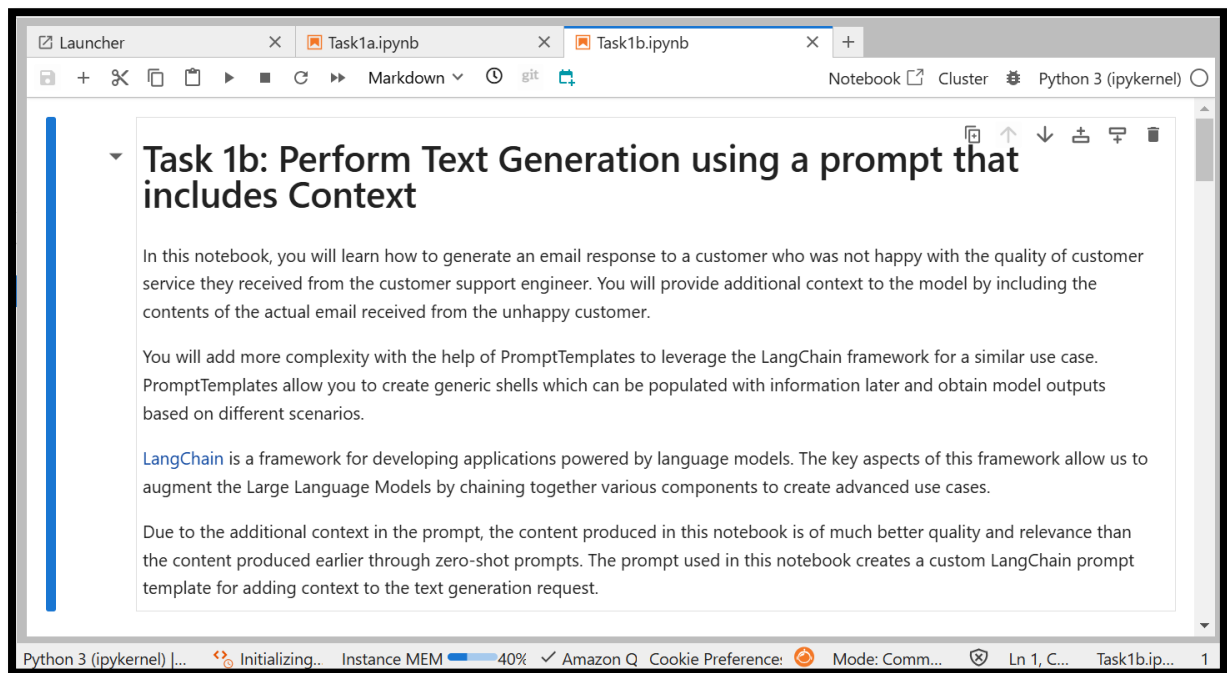
```
[6]: # invoke_model_with_response_stream
output = []
try:
    response = bedrock_client.invoke_model_with_response_stream(body=body, modelId=modelId, accept=accept, contentType=contentType)
    stream = response.get('body')

    i = 1
    if stream:
        for event in stream:
            chunk = event.get('chunk')
            if chunk:
                chunk_obj = json.loads(chunk.get('bytes').decode())
                text = chunk_obj['outputText']
                output.append(text)
```





Task1b.ipynb, which used the LangChain framework to communicate with the Amazon Bedrock API and created a custom LangChain prompt template to add context to the text generation request.



The screenshot shows a Jupyter Notebook interface with a tab labeled 'Task1b.ipynb'. The notebook content includes a title, an introductory paragraph about generating email responses, a paragraph about using PromptTemplates, a paragraph about the LangChain framework, and a concluding paragraph about the quality of the generated content.

## Task 1b: Perform Text Generation using a prompt that includes Context

In this notebook, you will learn how to generate an email response to a customer who was not happy with the quality of customer service they received from the customer support engineer. You will provide additional context to the model by including the contents of the actual email received from the unhappy customer.

You will add more complexity with the help of PromptTemplates to leverage the LangChain framework for a similar use case. PromptTemplates allow you to create generic shells which can be populated with information later and obtain model outputs based on different scenarios.

LangChain is a framework for developing applications powered by language models. The key aspects of this framework allow us to augment the Large Language Models by chaining together various components to create advanced use cases.

Due to the additional context in the prompt, the content produced in this notebook is of much better quality and relevance than the content produced earlier through zero-shot prompts. The prompt used in this notebook creates a custom LangChain prompt template for adding context to the text generation request.



The screenshot shows a Jupyter Notebook interface with a tab labeled 'Task1b.ipynb'. The notebook content includes a title, an introductory paragraph about setting up the environment, and a code cell with Python code for importing libraries and setting up the environment.

## Task 1b.1: Environment setup

In this task, you set up your environment.

```
[1]: #Create a service client by name using the default session.
import json
import os
import sys
import warnings

import boto3

warnings.filterwarnings('ignore')
module_path = ".."
sys.path.append(os.path.abspath(module_path))

bedrock_client = boto3.client('bedrock-runtime', region_name=os.environ.get("AWS_DEFAULT_REGION", None))
```

LauncherTask1a.ipynbTask1b.ipynbNotebookClusterPython 3 (ipykernel)

## Task 1b.2: Invoke the Bedrock LLM Model

In this task, you create an instance of the Bedrock class from `llms`. This expects a `model_id` which is the Amazon Resource Name (ARN) of the model available in Amazon Bedrock.

Optionally, you can pass a previously created boto3 client as well as some `model_kwargs` which can hold parameters such as `temperature`, `top_p`, `max_token_count`, or `stop_sequences` (more information on parameters can be explored in the Amazon Bedrock console).

Refer to [documentation](#) for Available text generation model Ids under Amazon Bedrock.

**Note:** The different models support different `model_kwargs`.

```
[2]: # Model configuration
from langchain_aws import ChatBedrock
from langchain_core.output_parsers import StrOutputParser

model_id = "meta.llama3-8b-instruct-v1:0"
model_kwargs = {
    "max_gen_len": 512,
    "temperature": 0,
    "top_p": 1,
}

# LangChain class for chat
chat_model = ChatBedrock(
    client=bedrock_client,
    model_id=model_id,
    model_kwargs=model_kwargs,
)
```

LauncherTask1a.ipynbTask1b.ipynbCodePython 3 (ipykernel)

## Task 1b.3: Create a LangChain custom prompt template

In this task, you will create a template for the prompt that you can pass different input variables on every run. This is useful when you have to generate content with different input variables that you may be fetching from a database.

In the previous task, we hardcoded the prompt. It might be the case that you have multiple customers sending similar negative feedback, and you now want to use each of those customers' emails and respond to them with an apology, but you also want to keep the response a bit personalized. In the following cell, you will explore how you can create a `PromptTemplate` to achieve this pattern.

```
[3]: # Create a prompt template that has multiple input variables
from langchain.prompts import PromptTemplate

multi_var_prompt = PromptTemplate(
    input_variables=["customerServiceManager", "customerName", "feedbackFromCustomer"],
    template="""
Human: Create an apology email from the Service Manager {customerServiceManager} at AnyCompany to {customerName} in
<customer_feedback>
{feedbackFromCustomer}
</customer_feedback>
    """
)
```

```
Assistant:"""
)

# Pass in values to the input variables
prompt = multi_var_prompt.format(customerServiceManager="Bob Smith",
                                customerName="John Doe",
                                feedbackFromCustomer="""Hello Bob,

I am very disappointed with the recent experience I had when I called your customer support.
I was expecting an immediate call back but it took three days for us to get a call back.
The first suggestion to fix the problem was incorrect. Ultimately the problem was fixed after three days.
We are very unhappy with the response provided and may consider taking our business elsewhere.
""")

[4]: # get number of tokens
num_tokens = chat_model.get_num_tokens(prompt)
print(f"Our prompt has {num_tokens} tokens")

tokenizer_config.json: 100% ██████████ 26.0/26.0 [00:00<00:00, 1.28kB/s]
vocab.json: 100% ██████████ 1.04M/1.04M [00:00<00:00, 43.8MB/s]
```

Python 3 (ipykernel) Initialized (additional servers n Instance MEM 4 ✓ Amazon Cookie Preferen Mode: Com... Ln 1, ... Task1b.i... 1

```
Launcher Task1a.ipynb Task1b.ipynb
+ - - - - -
Markdown Python 3 (ipykernel)

[5]: #invoke
response = chat_model.invoke(prompt)

[6]: # Configure a Chain to parse output
chain = StrOutputParser()
formatted_response=chain.invoke(response)
print(formatted_response)

Here is an apology email from Service Manager Bob Smith at AnyCompany to John Doe:

Subject: Apology and Follow-up on Recent Customer Support Experience

Dear John,

I am writing to express my sincerest apologies for the disappointing experience you had with our customer support team. I am Bob Smith, the Service Manager at AnyCompany, and I take full responsibility for the issues you faced.

I am deeply sorry that it took three days for us to get back to you after your initial call. This is not the level of service we strive to provide, and I understand how frustrating it must have been for you to wait for a resolution.

I am also sorry that the initial suggestion to fix the problem was incorrect. I can only imagine how frustrating that must have been for you, and I want to assure you that we are taking steps to ensure that our team is better equipped to provide accurate solutions moving forward.
```

I am also sorry that the initial suggestion to fix the problem was incorrect. I can only imagine how frustrating that must have been for you, and I want to assure you that we are taking steps to ensure that our team is better equipped to provide accurate solutions moving forward.

I am pleased to hear that the problem was ultimately resolved after three days, but I understand that the delay and incorrect solution caused unnecessary inconvenience and stress.

I want to assure you that we value your business and appreciate the opportunity to serve you. I would like to offer a gesture of goodwill to make up for the poor experience you had. Please let me know if there is anything I can do to make things right.

If you would like to discuss this further or have any feedback on how we can improve our service, please don't hesitate to reach out to me directly. Your feedback is invaluable in helping us to improve our customer experience.

Once again, I apologize for the poor service you received, and I hope to have the opportunity to serve you better in the future.

Thank you for your patience and understanding.

Sincerely,

Bob Smith  
Service Manager, AnyCompany